

JErlang: Erlang with Joins

Hubert Plociniczak

`hp105@doc.ic.ac.uk`

Supervisor: Professor Susan Eisenbach

Second marker: Professor Sophia Drossopoulou

June 22, 2009

Erlang

- Functional language
- Dynamic typing
- Concurrency Oriented Programming (Actors)
- Fault-tolerant, non-stop running
- Single variable assignment
- Powerful pattern-matching
- Open Telecom Platform (for example **gen_server** behaviour used for implementing the server in the client-server model)

Erlang - Processes and Communication

Create process

```
Pid = spawn(fun() -> io:format("New Process!", []))
```


Send message

```
Pid ! {john, salary, 22000}
```

Receive message

```
1 receive  
2     {ok, Value} ->  
3         process_value(Value);  
4     {error, Reason} ->  
5         shutdown(Reason)  
6 after  
7     Timeout -> process_no_response()  
8 end
```

How to synchronise on more than one message? (Wrong)



```
1  receive
2      {get, X} ->
3          receive
4              {set, Y} when (X == Y) ->
5                  {found, 2, X}
6          end;
7      {set, X} ->
8          receive
9              {get, Y} when (X == Y) ->
10                 {found, 2, X}
11          end
12  end
```

How to synchronise on more than one message? (Almost)

```
1 fun() ->
2   A = fun(ReceiveFunc) ->
3     receive
4       {get, X} ->
5         receive
6           {set, Y} when (X == Y) ->
7             {found, 2, X}
8         after 0 ->
9           self() ! {get, X},
10          ReceiveFunc(ReceiveFunc)
11       end
12   {set, X} ->
13     receive
14       {get, Y} when (X == Y) ->
15         {found, 2, X}
16     after 0 ->
17       self() ! {set, X},
18       ReceiveFunc(ReceiveFunc)
19     end
20   end
21 end,
22 A(A)
23 end.
```

Join - calculus

- Based on a Abstract Chemical Machine used for chemical reactions
- Process Calculi developed at INRIA
- Equivalent to Π -calculus
- Introduces multi-way join patterns
- Non-deterministic choice

Syntax

<i>function</i>	$::=$	$\overline{\text{functiondef}}$
<i>functiondef</i>	$ $	$(\bar{p}) \text{ when } guard \rightarrow e$
<i>expr(e)</i>	$::=$	$e_1 \text{ op}_e e_2$
	$ $	$e(\bar{e})$
	$ $	$\text{case } e \text{ of } \overline{\text{match}} \text{ end}$
	$ $	receive $\overline{\text{join}}$ end
	$ $	$e_1 ! e_2$
	$ $	e_1 , e
	$ $	$p =_m e_2$
	$ $	$\text{basicvalue} \mid \text{varld} \mid \{ \bar{e} \} \mid [\bar{e}]$
<i>basicvalue</i>	$::=$	$\text{atom} \mid \text{number} \mid \text{pid} \mid \text{funcl d}$
<i>value(v)</i>	$::=$	$\text{basicvalue} \mid \{ \bar{v} \} \mid [\bar{v}]$
<i>pattern(p)</i>	$::=$	$\text{varld} \mid \text{basicvalue} \mid \{ \bar{p} \} \mid [\bar{p}]$
<i>match</i>	$::=$	$p \text{ when } g \rightarrow e$
join	$::=$	$\text{jpattern} [\text{join_aux}] [\text{when } g] \rightarrow e$
join_aux	$::=$	and $\text{jpattern} [\text{join_aux}]$
jpattern	$::=$	$\text{propagation } p$
<i>propagation</i>	$::=$	true $ $ false
<i>guard(g)</i>	$::=$	$g_1 \text{ op}_g g_2 \mid \text{basicvalue}$
	$ $	$\text{varld} \mid g(\bar{g}) \mid \{ \bar{g} \} \mid [\bar{g}]$

Structural Operational Semantics

$$\text{Var}_0: \frac{F(\text{varId}) = v \quad v \neq \text{Udf}}{\text{varId}, F, Q \rightsquigarrow_E v, F, Q}$$

$$\text{Seq}: \frac{}{(v, e), F, Q \rightsquigarrow_E e, F, Q}$$

$$\text{Receive}_1: \frac{Q^a \leq_q Q \quad \forall (i \in 1..n) (join_i =_s jpattern_{i,1} \text{ and } \dots jpattern_{i,n} \text{ when } g_i \rightarrow e_i) \quad k \in 1..n, joinMatches(join_k, F, Q, Q^a, F', Q')}{\text{receive } join_1 \dots join_n \text{ end}, F, Q \rightsquigarrow_E e_k, F', Q'}$$

$$\text{Receive}_2: \frac{\forall (Q'' \leq_q Q). \forall (1 \leq l \leq n). \neg \exists F', Q' (joinMatches(join_l, F, Q'', F', Q'))}{\text{receive } join_1 \dots join_n \text{ end}, F, Q \rightsquigarrow_E \text{error}, F, Q}$$

Joins First Match Semantics

- There is a join which can be satisfied with the subset of the current mailbox
- For any previous join, with the subset of the mailbox given above, matching is unsuccessful
- For any smaller subset of the mailbox none of the defined joins can be satisfied

$$\begin{aligned}
 & Q^a \leq_q Q \\
 & \forall (i \in 1..n) \text{ (join}_i =_s \text{jpattern}_{i,1} \text{ and } \dots \text{jpattern}_{i,n} \text{ when } g_i \rightarrow e_i) \\
 & k \in 1..n, \text{ joinMatches (join}_k, F, Q, Q^a, F', Q') \\
 & \forall (Q^b <_q Q^a). \forall (1 \leq l \leq n). \neg \exists F'', Q'' \\
 & \quad (\text{joinMatches (join}_l, F, Q, Q^b, F'', Q'')) \\
 & \forall (1 \leq l < k). \neg \exists F'', Q'' \\
 & \quad (\text{joinMatches (join}_l, F, Q, Q^a, F'', Q''))
 \end{aligned}$$

Receive_{First-Match}: $\frac{\text{receive join}_1 \dots \text{join}_n \text{ end, } F, Q \rightsquigarrow_E e_k, F', Q'}{\text{receive join}_1 \dots \text{join}_n \text{ end, } F, Q \rightsquigarrow_E e_k, F', Q'}$

JErlang language features

- Joins
- Guards
- Timeouts
- Non-linear patterns
- Propagation
- Synchronous calls
- new OTP design pattern

Joins

- First Match semantics
- The possibility of having multiple patterns (possibly the same)
- Take into account the context (bounded vs unbounded variables)
- Messages' order preserved

```
1 operation() ->
2   receive
3     {ok, sum} and {val, X} and {val, Y} ->
4       {sum, X + Y};
5     {ok, mult} and {val, X} and {val, Y} ->
6       {mult, X * Y};
7     {ok, sub} and {val, X} and {val, Y} ->
8       {sub, X - Y}
9   end
10 end.
```

Guards

Provides additional filtering not expressible in terms of patterns.
Limited number of expressions without side-effects (Erlang constraint).

```
1 receive
2     {Transaction, M} and {limit, Lower, Upper}
3         when (Lower <= M and M <= Upper) ->
4         commit_transaction(M, Transaction)
5 end
```

Timeouts

- Not specified in Join-calculus
- Integral part of Erlang **receive** construct
- Time measured for *idle* periods

```
1 cash_machine_authorise(Timeout) ->
2   receive
3     {pin, Pin} and {card, Number} ->
4       authenticate(Pin, Number);
5     abort and {card, Number} ->
6       abort_authentication(Number)
7   after Timeout ->
8     take_card(Number, Timeout);
9   end
10 end.
```

Non-linear patterns

- Messages can match multiple joins
- Unbounded variables spanned over multiple patterns have to agree on values
- Computationally expensive

```
1 receive
2   {get, X} and {set, X} ->
3     {found, 2, X}
4 end
5 ...
6 receive
7   {Pin, Id} and {auth, Pin} and {commit, Id} ->
8     perform_transaction(Pin, Id)
9 end.
```

Propagation

Inspired by Constraint Handling Rules and Haskell Join Rules language.

Allows for copying correct messages instead of removing them.

```
1 receive  
2   prop({session , Id}) and {act , Action , Id} ->  
3     perform_action(Action , Id);  
4   {session , Id} and {logout , Id} ->  
5     logout_user(Id)  
6 end
```

Synchronous calls

Explicit synchronisation requires process identifiers in the message.

```
1 receive  
2   {accept, Pid1} and {asynchronous, Value}  
3               and {accept, Pid2}    ->  
4       Pid1 ! {ok, Value},  
5       Pid2 ! {ok, Value}  
6 end
```


gen_joins

Open Telecom Platform design pattern (in comparison to **gen_server** allows for synchronisation on multiple messages)

```
1  get() ->
2      jerlang_gen_joins:call(?MODULE, get).
3
4  set(Value) ->
5      jerlang_gen_joins:cast(?MODULE, {set, Value}).
6
7  %% CALLBACKS
8  handle_join( {set, Value} and get, Status) ->
9      {[noreply, {reply, {ok, Value}}],
10       [Value | Status]}.
```

Definition (Santa Claus problem)

Santa Claus sleeps at the North pole until awakened by either all of the nine reindeer, or by a group of three out of ten elves. He performs one of two indivisible actions:

- If awakened by the group of reindeer, Santa harnesses them to a sleigh, delivers toys, and finally unharnesses the reindeer who then go on holiday.
- If awakened by a group of elves, Santa shows them into his office, consults with them on toy R&D, and finally shows them out so they can return to work constructing toys.

A waiting group of reindeer must be served by Santa before a waiting group of elves.

Santa Claus Demo

Santa Claus solution in `gen_joins`

Santa Claus - Minimal Erlang solution

```
1 worker(Secretary, Message) ->
2   receive after random:uniform(1000) -> ok end, % random delay
3   Secretary ! self(), % send my PID to the secretary
4   Gate_Keeper = receive X -> X end, % await permission to enter
5   io:put_chars(Message), % do my action
6   Gate_Keeper ! {leave, self()}, % tell the gate-keeper I'm done
7   worker(Secretary, Message). % do it all again
8
9 secretary(Santa, Species, Count) ->
10  secretary_loop(Count, [], {Santa, Species, Count}).
11
12 secretary_loop(0, Group, {Santa, Species, Count}) ->
13  Santa ! {Species, Group},
14  secretary(Santa, Species, Count);
15 secretary_loop(N, Group, State) ->
16  receive PID ->
17    secretary_loop(N-1, [PID|Group], State)
18  end.
19
20 santa() ->
21  {Species, Group} =
22  receive % first pick up a reindeer group
23    {reindeer, G} -> {reindeer, G} % if there is one, otherwise
24  after 0 ->
25    receive % wait for reindeer or elves,
26      {reindeer, G} -> {reindeer, G}
27    ; {elves, G} -> {elves, G}
28    end % whichever turns up first.
29  end,
30  case Species
31  of reindeer -> io:format("Ho, ho, ho! Let's deliver toys!~n", [])
32  ; elf -> io:format("Ho, ho, ho! Let's meet in the study!~n", [])
33  end,
34  [PID ! self() || PID <- Group], % tell them all to enter
35  [receive {leave, PID} -> ok end % wait for each of them to leave
36  || PID <- Group],
37  santa().
38
39 spawn_worker(Secretary, Before, I, After) ->
40  Message = Before ++ integer_to_list(I) ++ After,
41  spawn(fun () -> worker(Secretary, Message) end).
```

Santa Claus - Minimal JErlang solution


```
1  santa() ->
2      io:format("It was a long night. Time to bed~n"),
3      Group =
4      receive
5          {reindeer, Pid1} and {reindeer, Pid2} and
6          {reindeer, Pid3} and {reindeer, Pid4} and
7          {reindeer, Pid5} and {reindeer, Pid6} and
8          {reindeer, Pid7} and {reindeer, Pid8} and
9          {reindeer, Pid9} ->
10         io:format("Ho, ho, ho! Let's deliver presents!~n"),
11         [Pid1, Pid2, Pid3, Pid4,
12          Pid5, Pid6, Pid7, Pid8, Pid9];
13         {elf, Pid1} and {elf, Pid2} and {elf, Pid3} ->
14         io:format("Ho, ho, ho! Let's discuss R&D possibillites!~n"),
15         [Pid1, Pid2, Pid3]
16     end,
17     [Pid ! ok || Pid <- Group],
18
19     %% Synchronise on return of the animals
20     receive
21         {r, done} and {r, done} and {r, done} and {r, done} and {r, done}
22         and {r, done} and r, done} and {r, done} and {r, done} ->
23         ok;
24         {e, done} and {e, done} and {e, done} ->
25         ok
26     end,
27     santa().
28
29 worker(Santa, Type, Id, Action, Ans) ->
30     worker1(Santa, Type, Id, Action, Ans).
31
32 worker1(Santa, Type, Id, Action, Ans) ->
33     receive after random:uniform(1000) -> ok end,
34     Santa ! {Type, self()},
35     io:format("~p ~p: Waiting at the gate~n", [Type, Id]),
36     receive ok -> ok end,
37     io:format("~p ~p: ~p~n", [Type, Id, Action]),
38     Santa ! {Ans, done},
39     worker1(Santa, Type, Id, Action, Ans).
```

Sad Santa Claus - Minimal JErlang solution

No additional “hacks” for proper prioritisation and clear solution for different types of messages

```
1 receive
2     {reindeer, Pid1} and {reindeer, Pid2} and {reindeer, Pid3}
3     and {reindeer, Pid4} and {reindeer, Pid5} and {reindeer, Pid6}
4     and {reindeer, Pid7} and {reindeer, Pid8} and {reindeer, Pid9} ->
5         io:format("Ho,ho,ho! Let's deliver presents!~n"),
6         [Pid1, Pid2, Pid3, Pid4,
7          Pid5, Pid6, Pid7, Pid8, Pid9];
8     {ork,sergant,Pid1} and {ork,captain,Pid2} and {ork,sergant,Pid3} ->
9         io:format("Ho,ho,ho? No presents, orks destroyed the factory!~n"),
10        [Pid1, Pid2, Pid3];
11     {elf, Pid1} and {elf, Pid2} and {elf, Pid3} ->
12         io:format("Ho,ho, ho!Let's discuss R&D possibillites!~n"),
13         [Pid1, Pid2, Pid3]
14 end
```

General architecture



parse_transform and Abstract Syntax Trees

- Different semantics to the valid Erlang syntax code
- Requires only a single line of code in the original module

```
-compile({parse_transform, jerlang_parse}).
```

- Runs as part of the parser stage in the compiler
- JErLang provides various version of **parse_transform** for different modes
- Independent of the Erlang release, compiler, run-time and the application
- Uses familiar Erlang syntax for analysis of the code (pattern matching, tail recursion)
- Enables pretty joins syntax instead of complex calls to JErLang's library
- Required careful analysis of the code to avoid spurious warnings and errors
- “Not recommended” for standard Erlang programmers

Santa Claus (**gen_joins**) without parse_transform

Almost 400 lines of code of condensed code

[illegible][illegible][illegible]

Santa Claus (**gen_joins**) with parse_transform

67 lines of readable code

```
1 module [xdebug_gen_joins_test_santa_parse].
2 compile [{ parse_transform, xdebug_gen_joins_parse }].
3
4 ~behaviour [xdebug_gen_joins].
5
6 ~export [init/1, handle_join/2, terminate/1].
7 ~export [start/1, stop/1, elf/1, window/1].
8 ~export [elf_done/1, window_done/1, status/1].
9
10 start [] ->
11   xdebug_gen_joins:start [{ global, ?MODULE }, ?MODULE, [], []].
12
13 stop [] ->
14   xdebug_gen_joins:call [{ global, ?MODULE }, stop].
15
16 terminate [] ->
17   ok.
18
19 init [] ->
20   {ok, {{1,1}, sleeping}}.
21
22 elf [] ->
23   xdebug_gen_joins:call [{ global, ?MODULE }, elf, infinity],
24   ok.
25
26 elf_done [] ->
27   xdebug_gen_joins:call [{ global, ?MODULE }, {done, elf}].
28
29 window [] ->
30   xdebug_gen_joins:call [{ global, ?MODULE }, window, infinity].
31
32 window_done [] ->
33   xdebug_gen_joins:call [{ global, ?MODULE }, {done, window}].
34
35 status [Status] ->
36   xdebug_gen_joins:call [{ global, ?MODULE }, {status, Status}, infinity].
37
38 handle_join [{status, {A, B}}, {{C, D}, _}=S] when ||A<=C|| and ||B<=D||
39   ->
40     io:format("I am done ~s", [{C, D}]),
41     [{ {reply, ok}, _ }, S],
42     handle_join(stop, _) ->
43       io:format("Stopping the server~s", []),
44       {stop, normal};
45     handle_join [{done, window} and {done, window} and {done, window} and
46       {done, window} and {done, window} and {done, window} and
47       {done, window} and {done, window} and {done, window}],
48       {Counter, awake_window} ->
49       io:format("All windows returned~s", []),
50       [{ {reply, []} _ <- lists:seq(1,1)}, {Counter, sleeping}];
51
52     handle_join [{done, elf} and {done, elf} and {done, elf}],
53       {Counter, awake_elf} ->
54       io:format("All elves returned~s", []),
55       [{ {reply, []} _ <- lists:seq(1,1)}, {Counter, sleeping}];
56
57     handle_join [window and window and window and
58       window and window and window and
59       window and window and window,
60       {{Window, {Ehon}}, sleeping}] ->
61       io:format("Ho, ho, ho! Let's all [yay] presents~s", []).
```

VM or non-VM

Self-contained JErland library:


- Portable
- Requires only a single line in the original code
- Introduces overhead related to mailbox access/storage (library mailbox with processed messages and the unprocessed ones in the VM)

VM-accelerated JErland library:

- Implemented without any documentation of VM or compiler (C, Erlang)
- Requires patched version of the Erlang compiler and run-time
- Provides primitives necessary to efficiently access the mailbox
- Uses modified version of the self-contained library
- Significantly faster in many situations
- Uses hash-map data structure to complement VM's queue

RETE algorithm

- Well used in Production Rule Systems for more than 30 years
- Initially applicable to **gen_joins** implementation, used in VM and non-VM implementation
- Avoids re-computation of results (better performance for larger mailboxes and complicated patterns)




Alpha memory

Beta memory

Performance

- Roughly similar when comparing respective language features in other implementations
- Suffers from “too-much-mail is bad for you” problem
- Various optimisation techniques used



Further optimisations

- Joins' patterns ordering

```
1 receive
2     {foo, One} and test1 and test2 and {bar, One} ->
3     ...
4 end
```

VS

```
1 receive
2     {foo, One} and {bar, One} and test1 and test2 ->
3     ...
4 end
```

- Early application of variables for faster message filtering
- Detection of identical messages

Future work

- Parallel joins solver (constraints on the expressiveness)
- Various static analysis of the code patterns to boost performance
- More efficient VM-implementation (data structures)
- Formalisation of the equivalence of Erlang and JErlang
- Go open-source (roughly 7 000 lines of code without VM changes)

Achievements

- Defined formal syntax and semantics of the extension
- Implemented extension of Erlang with consistent semantics that increases the expresiveness of the language and provides features that do not exist in other joins implementations
- Implemented self-standing JErland library with and without VM/compiler changes
- Used static analysis techniques for producing better code
- Investigated various optimisation techniques that significantly increase the usability of joins

Santa Claus problem - Polyphonic C# (C ω)

Elegant? Intuitive?

```
1 public class nway {
2     public async produce(int n) & public void consume() {
3         if (n==1) {
4             alldone();
5         } else {
6             produce(n-1);
7         }
8     }
9
10    public void waitforalldone() & async alldone() {
11        return;
12    }
13 }
```

Santa Claus problem - Polyphonic C# part 2

```
1  class santa {
2      static nway harness = new nway();
3      static nway unharness = new nway();
4      static nway roomin = new nway();
5      static nway roomout = new nway();
6
7      static void santalife() {
8          while (true) {
9              waittobewoken();
10             // get back here after dealing with elves or reindeer
11         }
12     }
13
14     static void waittobewoken() & static async elvesready() {
15         roomin.produce(3);
16         roomin.waitforalldone();
17         elveswaiting(0);
18         // all elves in the room, consult
19         roomout.produce(3);
20         roomout.waitforalldone();
21         // all elves shown out, go back to bed
22     }
23
24     static void waittobewoken() & static async reindeerready() {
25         // similar to elvesready chord
26     }
```

Santa Claus problem - Polyphonic C# part 3

```
1 static async elflife(int elfid) {
2     while (true) {
3         // work
4         elfqueue();           // wait to join group of 3
5         roomin.consume();     // wait to be shown in
6         // consult with santa
7         roomout.consume();    // wait to be shown out again
8     }
9 }
10
11 static void elfqueue() & static async elveswaiting(int e) {
12     if (e==2) {
13         elvesready();        // wake up santa
14     }
15     else {
16         elveswaiting(e+1);
17     }
18 }
```

Expandable? No prioritisation - requires further, hard to maintain, “hacks”.

JErlang allows for much more intuitive solution of the problem.

NataMQ

